

Assignment - Matheus Douglas MATOS MARCONDES

November 25, 2019

```
In [1]: import numpy as np
import numpy.random as rand
import matplotlib.pyplot as plt

from sklearn.datasets import make_blobs, make_moons, make_circles
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import OneHotEncoder
```

1 Perceptrons and Back propagation

2 Submit to moodle by the 25th November - Please submit a PDF version of your notebook

2.1 One layer perceptron - Linear separation


We start by generating *blobs*. In this setting data is linearly separable, we can thus use a single layer perceptron.

$$\hat{y} = W^T x + b$$

$$\hat{P}(y_j = 1) = \text{Softmax}(\hat{y})_j$$

With $W \in M_{P,N}(\mathbb{R})$ and b a vector of size N . P, N the dimension of the input and the number of classes.

```
In [2]: X, Y = make_blobs(n_samples=500, n_features=2)
plt.figure(figsize=(5,5))
plt.scatter(X[:,0], X[:,1], c=Y)
plt.show()
Y = OneHotEncoder().fit_transform(Y.reshape(-1,1)).toarray()
```



```
In [3]: X_train, X_test, Y_train, Y_test = train_test_split(X, Y, random_state=42,
                                                             test_size=0.4)
```

2.1.1 Task 1: Implement softmax and the forward pass of a single layer perceptron.

```
In [4]: def softmax(y):
        """
        Takes the output of a layer as input and returns a probability distribution
        input:
            y (np.array)

        returns:
            y (np.array)
        """
        aux = y.reshape(-1)
        aux = np.exp(aux)
        somme = aux.sum()
        aux = aux / somme

        return aux.reshape(-1,1)

def forward_one_layer(W, b, x):
```

```

"""
Computes the forward pass of a single layer perceptron
input:
    W (np.array): (INPUT_SHAPE, N_CLASSES) The weight matrix of the perceptron
    b (np.array): (N_CLASSES, 1) The bias matrix of the perceptron
    x (np.array): (INPUT_SHAPE, 1) The input of the perceptron

returns:
    (np.array) (N_CLASSES, 1)
"""
return W.T @ x + b

```

The loss typically associated with a classification problem is the cross entropy loss:

$$loss = -\log(\hat{P}(y_j = 1))$$

2.1.2 Question 1 (2 points): Derive the gradients of loss with respect to W and b .

$$\nabla_W loss = ?$$

$$\nabla_b loss = ?$$

Answer: We can write the loss function as:

$$\begin{aligned}
 loss &= -\log(\hat{P}(y_j = 1)) \\
 &= -\log(\text{Softmax}(\hat{y})_j) \\
 &= -\log\left(\frac{e^{y_j}}{\sum_{k=1}^C e^{y_k}}\right) \\
 &= \log\left(\sum_{k=1}^C e^{y_k}\right) - \log e^{y_j} \\
 &= \log\left(\sum_{k=1}^C e^{y_k}\right) - y_j
 \end{aligned}$$

where $y_k = W_{(k,\cdot)}^T x + b_k = \sum_{i=1}^C w_{i,k} x_i + b_k$.

Then, we have:

- For $k \neq j$:

$$\begin{aligned}
 \frac{\partial loss}{\partial W_{i,k}} &= \frac{e^{y_k}}{\sum_{p=1}^C e^{y_p}} x_i \\
 &= \text{Softmax}(\hat{y})_k x_i
 \end{aligned}$$

$$\begin{aligned}
 \frac{\partial loss}{\partial b_k} &= \frac{e^{y_k}}{\sum_{p=1}^C e^{y_p}} \\
 &= \text{Softmax}(\hat{y})_k
 \end{aligned}$$

- For $k = j$:

$$\begin{aligned}
 \frac{\partial loss}{\partial W_{i,j}} &= \frac{e^{y_j}}{\sum_{p=1}^C e^{y_p}} x_i - x_i \\
 &= x_i (\text{Softmax}(\hat{y})_j - 1)
 \end{aligned}$$

$$\begin{aligned}\frac{\partial \text{loss}}{\partial b_j} &= \frac{e^{y_j}}{\sum_{p=1}^C e^{y_p}} - 1 \\ &= \text{Softmax}(\hat{y})_j - 1\end{aligned}$$

Hence, we have:

$$\begin{aligned}\nabla_{W\text{loss}} &= \begin{bmatrix} x_1 \\ \vdots \\ x_N \end{bmatrix} ([\text{Softmax}(\hat{y})_1 \quad \cdots \quad \text{Softmax}(\hat{y})_C] - [0 \quad \cdots \quad 0 \quad 1_{jth} \quad 0 \quad \cdots]) \\ &= x (\text{Softmax}(\hat{y})^T - y^T)\end{aligned}$$

$$\begin{aligned}\nabla_{b\text{loss}} &= \begin{bmatrix} \text{Softmax}(\hat{y})_1 \\ \vdots \\ \text{Softmax}(\hat{y})_C \end{bmatrix} - \begin{bmatrix} 0 \\ \vdots \\ 0 \\ 1_{jth} \\ 0 \\ \vdots \end{bmatrix} \\ &= \text{Softmax}(\hat{y}) - y\end{aligned}$$

2.1.3 Task 2: Implement the gradients

```
In [5]: def compute_grads_one_layer(softmaxed, x, y):
        """
        inputs:
            softmaxed (np.array): (N_CLASSES, 1)
            y (np.array): (N_CLASSES, 1)
            x (np.array): (INPUT_SHAPE, 1)

        returns:
            d_W (np.array): (INPUT_SHAPE, N_CLASSES) Gradient of the loss with
                           respect to the weight matrix
            d_b (np.array): (N_CLASSES, 1) Gradient of the loss with respect to
                           the bias matrix
        """
        d_W = x @ (softmaxed - y).T
        d_b = softmaxed - y

        return d_W, d_b

def compute_loss(softmaxed, y):
    """
    inputs:
        softmaxed (np.array): (N_CLASSES, 1)
        y (np.array): (N_CLASSES, 1)

    returns:
```

```

        (float)
    """
    return float( -np.log( softmaxed.T @ y ) )

```

2.1.4 Question 2 (1 points): As a sanity check, we want to compare the gradients we calculated to approximated gradients. How could we do this ?

Answer: In order to check if the gradients are corrects, we can calculate approximate gradients and compare the gradients to them. The approximate gradients can be calculated according to the partial derivative definition:

$$\frac{\partial f(x,y)}{\partial x} = \lim_{h \rightarrow 0} \frac{f(x+h,y) - f(x,y)}{h}$$

We calculate then for each element of the gradient of the function.

Example: For the *loss* with respect to *W*, we have:

$$\nabla_W \text{loss}_{(i,j)} = \lim_{h \rightarrow 0} \frac{\text{loss}(W + H) - \text{loss}(W)}{h}$$

where $H = (h_{m,n})$, with $h_{m,n} = \begin{cases} 0, & (m,n) \neq (i,j) \\ h, & (m,n) = (i,j) \end{cases}$.

2.1.5 Task 3: Implement the approx gradient function for the weight matrix

```

In [6]: def approx_grad_W(W, b, x, y, h=0.0001):
    """
    Approximates the gradient with respect to W
    input:
        W (np.array): (INPUT_SHAPE, N_CLASSES) The weight matrix of the perceptron
        b (np.array): (N_CLASSES, 1) The bias matrix of the perceptron
        x (np.array): (INPUT_SHAPE, 1) The input of the perceptron
        y (np.array): (N_CLASSES, 1) The input of the perceptron
        h (float): variation

    returns:
        d_W_approx (np.array): (INPUT_SHAPE, N_CLASSES)
    """
    # The approx gradient is (loss(W_(i,j)+h, b) - loss(W_(i,j), b)) / h
    y_0 = W.T @ x + b
    softmaxed = softmax(y_0)

    result = np.zeros(W.shape)
    for i in range(W.shape[0]):
        for j in range(W.shape[1]):
            h_matrix = np.zeros(W.shape)
            h_matrix[i][j] = h
            W_h = W + h_matrix
            y_0_h = W_h.T @ x + b
            softmaxed_h = softmax(y_0_h)

```

```

        result[i][j] = (
            compute_loss(softmaxed_h, y) - compute_loss(softmaxed, y) ) / h

    return result

```

The following function trains the perceptron

```

In [7]: def train_one_layer(X_train, Y_train, X_test, Y_test, lr,
                             n_it=1000, test_freq=10, random_seed=42):

    INPUT_SHAPE = X_train.shape[1]
    N_CLASSES = Y_train.shape[1]

    # Initialise metrics lists
    loss = []
    acc = []
    test_acc = []
    approx_grads = []
    grads = []

    # Initialisation of the weights
    np.random.seed(random_seed)
    b = rand.normal(size=(N_CLASSES, 1))
    W = rand.normal(size=(INPUT_SHAPE, N_CLASSES))

    # Shuffling data
    indexes = rand.randint(X_train.shape[0], size=n_it)
    # training loop
    for it, i in enumerate(indexes):
        x = X_train[i,:].reshape(-1,1)
        y = Y_train[i,:].reshape(-1,1)

        # Forward passs
        softmaxed = forward_one_layer(W, b, x)
        # Back propagation
        d_W, d_b = compute_grads_one_layer(softmaxed, x, y)
        W -= lr * d_W
        b -= lr * d_b

        # Recording approximate gradients
        grads.append(d_W)
        approx_grads.append(approx_grad_W(W, b, x, y))

        # Metrics recording
        loss.append(compute_loss(softmaxed, y))
        acc.append(np.argmax(softmaxed) == np.argmax(y))

    # Test loop

```

```

        if it % test_freq == 0:
            acc_temp = []
            for i in range(X_test.shape[0]):
                x = X_train[i,:].reshape(-1,1)
                y = Y_train[i,:].reshape(-1,1)
                softmaxed = forward_one_layer(W, b, x)
                acc_temp.append(np.argmax(softmaxed) == np.argmax(y))

            test_acc.append(np.mean(acc_temp))

    return W, b, loss, acc, test_acc, np.stack(grads, -1), np.stack(approx_grads, -1)

```

2.1.6 Task 4 (4 points): Using the previous function train the model. By producing an appropriate set of plots validate that it trained correctly and that you computed the correct gradients. Remember that a plot should be self explanatory.

In [8]: `W, b, loss, acc, test_acc, grads, approx_grads = train_one_layer(X_train, Y_train, X_test, Y_test, 0.01)`

```

iteration = [t+1 for t in range(1000)]

# Comparing the values between grads and approx_grads
error = []
for it in range(grads.shape[2]):
    aux = 0;
    for i in range(grads.shape[0]):
        for j in range(grads.shape[1]):
            aux += (grads[i][j][it] - approx_grads[i][j][it])**2
    error.append(aux)

plt.plot(iteration, error)
plt.xlabel("Iteration")
plt.ylabel("Error")
plt.title("Error between the grads and approx_grads for each iteration")
plt.show()

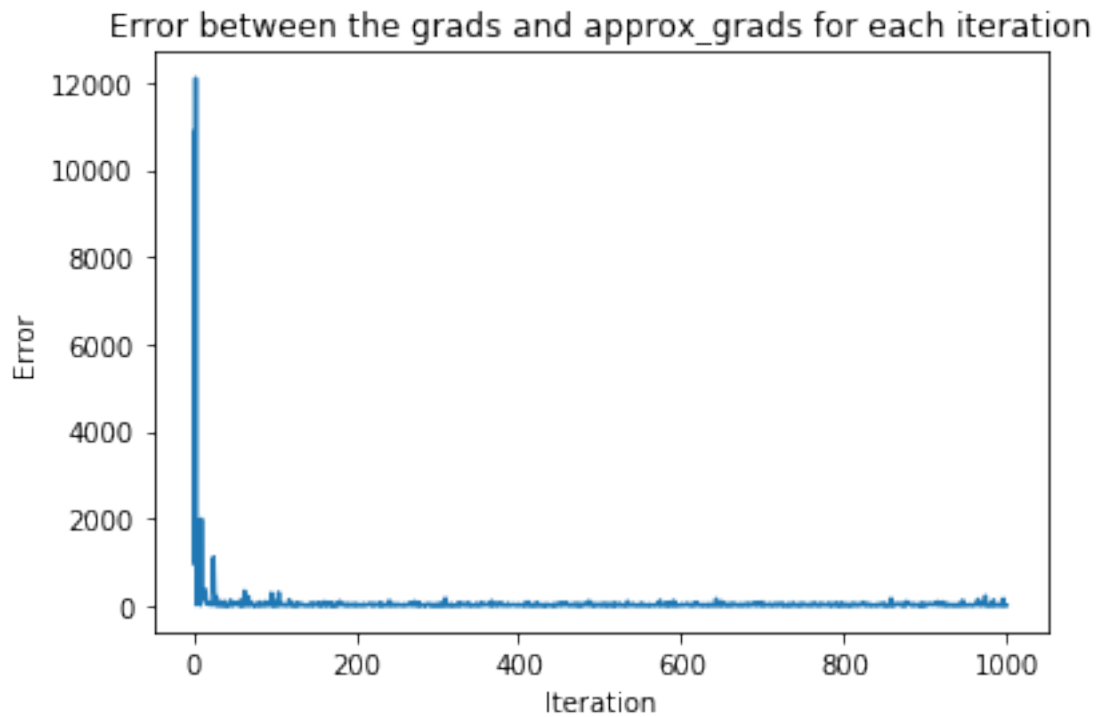
# evolution of loss function
plt.plot(iteration, loss)
plt.xlabel("Iteration")
plt.ylabel("Loss")
plt.title("Evolution of the Loss function")
plt.show()

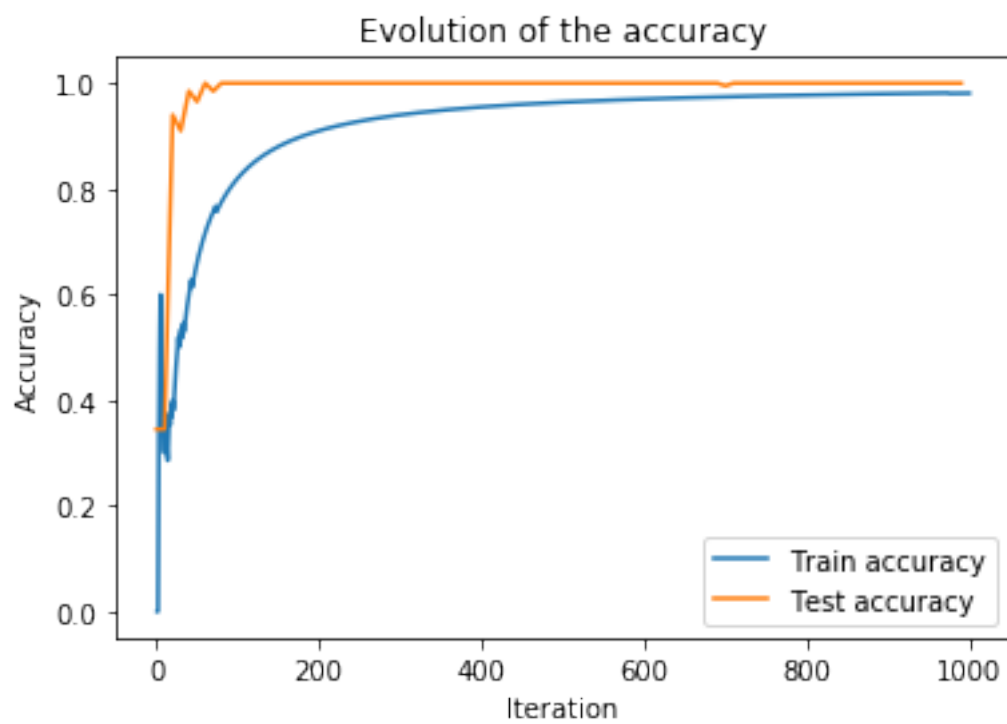

# Training data accuracy
freq = 10
acc_true = [a/b for a,b in zip(np.array(acc).cumsum(), range(1,len(acc)+1))]
plt.plot(iteration, acc_true, label="Train accuracy")
plt.plot([t*freq for t in range(100)], test_acc, label="Test accuracy")

```

```
plt.xlabel("Iteration")
plt.ylabel("Accuracy")
plt.title("Evolution of the accuracy")
plt.legend()
plt.show()
```

C:\Users\Matheus Douglas\Anaconda3\lib\site-packages\ipykernel_launcher.py:26: RuntimeWarning:





3 Non linearity

Now that the perceptron is trained, we can visualize its decision function.

```
In [9]: def plot_decision(X, Y, forward, figure=None):
        """Plots the decision function of a perceptron with respect to a
            forward function

            input:
                X, Y (np.array): Test data
                forward (function): only accepts x as input
                                    (Ex: lambda x: forward_one_layer(W, b, x))
                figure (plt.figure): optional, usefull if you dont want to generate
                                    any new figure, in the case of suplots.
            """
        markers=[".", "*", "D"]
        low0, high0 = np.min(X[:,0]), np.max(X[:,0])
        low1, high1 = np.min(X[:,1]), np.max(X[:,1])
        data = np.zeros((100,100,Y.shape[1]))
        for i1, x1 in enumerate(np.linspace(low0,high0,100)):
            for i2, x2 in enumerate(np.linspace(low1,high1,100)):
                x = np.array([x1, x2]).reshape(-1, 1)
                softmaxed = forward(x)
                data[i2, i1, :] = softmaxed.reshape(-1)
        if Y.shape[1] < 3:
            data = data[:, :, 0]


        if figure is None:
            plt.figure(figsize=(10,10))

        plt.imshow(data, extent=(low0,high0,low1,high1), origin='lower',
                    interpolation='gaussian')
        for c in range(Y.shape[1]):
            plt.scatter(X[np.argmax(Y, 1) == c, 0], X[np.argmax(Y, 1) == c, 1], c='k',
                        marker=markers[c], label="class %i" % c)

        plt.legend()
        plt.show()
```

```
In [10]: plot_decision(X_test, Y_test, lambda x: forward_one_layer(W, b, x))
```

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255]




The data we used was linearly separable, but what if it is not ?

3.0.1 Task 5 (1 point): Using the following data train a new perceptron and visualize its decision function.


```
In [11]: X, Y = make_moons(n_samples=500, noise=.2)
plt.figure(figsize=(10,10))
plt.scatter(X[:,0], X[:,1], c=Y)
plt.show()
Y = OneHotEncoder().fit_transform(Y.reshape(-1,1)).toarray()

X_train, X_test, Y_train, Y_test = train_test_split(X, Y, random_state=42,
                                                    test_size=0.4)
```



```
In [12]: W, b, loss, acc, test_acc, grads, approx_grads = train_one_layer(X_train, Y_train,
                                                                           X_test, Y_test, 0.01)
plot_decision(X_test, Y_test, lambda x: forward_one_layer(W, b, x))
```

C:\Users\Matheus Douglas\Anaconda3\lib\site-packages\ipykernel_launcher.py:26: RuntimeWarning:



4 Two Layer perceptron - Feed forward neural network

In neural networks, non linearity comes from having at least one hidden layer and a non linear activation function such as the *ReLU* function.

$$h = \text{ReLU}(W_h^\top x + b_h)$$

$$\hat{y} = W_o^\top h + b_o$$

$$\hat{P}(y_j = 1) = \text{Softmax}(\hat{y})_j$$

With $W_h \in M_{N,H}(\mathbb{R})$ and $W_o \in M_{C,H}(\mathbb{R})$ and b^h, b^o vectors of corresponding dimensions. H is the number of hidden units.

The *ReLU* function is defined as:

$$\text{ReLU}(x) = \begin{cases} 0 & \text{if } x < 0 \\ x & \text{otherwise} \end{cases}$$

4.0.1 Question 3 (2 points): Derive the gradients with respect to W_h and b_h .

Answer:

We have that:

$$\begin{aligned}
ReLU(x) &= \max(0, x) \\
ReLU'(x) &= \frac{dReLU}{dx}(x) \\
&= \mathbb{1}_{\{x>0\}} \\
&= \mathbb{1}_{\{ReLU(x)>0\}} \\
&= ReLU'(ReLU(x))
\end{aligned}$$

Also,

$$\begin{aligned}
h_i &= ReLU\left(W_{h(i,\cdot)}^T x + b_{h(i)}\right) \\
&= ReLU\left(\sum_k w_{h(k,i)} x + b_{h(i)}\right) \\
\frac{\partial h_i}{\partial W_{h(u,v)}} &= \begin{cases} 0 & , \quad i \neq v \\ x_u ReLU'\left(\sum_k w_{h(k,i)} x_k + b_{h(i)}\right) & , \quad i = v \end{cases} \\
&= \begin{cases} 0 & , \quad i \neq v \\ x_u ReLU'(h_i) & , \quad i = v \end{cases} \\
\frac{\partial h_i}{\partial b_{h(v)}} &= \begin{cases} 0 & , \quad i \neq v \\ ReLU'(h_i) & , \quad i = v \end{cases}
\end{aligned}$$

Then, for the \hat{y} , we have:

$$\begin{aligned}
\hat{y}_i &= W_{0(i,\cdot)}^T h + b_{0(i)} \\
&= \sum_k w_{0(k,i)} h_k + b_{0(i)} \\
\frac{\partial \hat{y}_i}{\partial W_{h(u,v)}} &= \sum_k w_{0(k,i)} \frac{\partial h_k}{\partial W_{h(u,v)}} \\
&= w_{0(v,i)} \frac{\partial h_v}{\partial W_{h(u,v)}} \\
&= w_{0(v,i)} x_u ReLU'(h_v) \\
\frac{\partial \hat{y}_i}{\partial b_{h(v)}} &= w_{0(v,i)} ReLU'(h_v)
\end{aligned}$$

Therefore, the gradients of the *loss* with respect to W_h and b_h are:

$$\text{loss} = \log(\sum_k e^{y_k}) - y_j$$

$$\begin{aligned} \frac{\partial \text{loss}}{\partial W_{h(u,v)}} &= \frac{\sum_k e^{y_k} \frac{\partial y_k}{\partial W_{h(u,v)}}}{\sum_k e^{y_k}} - \frac{\partial y_j}{\partial W_{h(u,v)}} \\ &= \sum_k \text{Softmax}(\hat{y})_k \frac{\partial y_k}{\partial W_{h(u,v)}} - \frac{\partial y_j}{\partial W_{h(u,v)}} \\ &= \sum_k \text{Softmax}(\hat{y})_k w_{0(v,k)} x_u \text{ReLU}'(h_v) - \\ &\quad w_{0(v,j)} x_u \text{ReLU}'(h_v) \\ &= x_u \text{ReLU}'(h_v) W_{0(v,\cdot)} [\text{Softmax}(\hat{y}) - y] \\ \frac{\partial \text{loss}}{\partial b_{h(v)}} &= \text{ReLU}'(h_v) W_{0(v,\cdot)} [\text{Softmax}(\hat{y}) - y] \end{aligned}$$

We can note that:

$$\nabla_{W_h} \text{loss} = x \nabla_{b_h} \text{loss}^T = \begin{bmatrix} x_1 \\ \vdots \\ x_p \end{bmatrix} \begin{bmatrix} \nabla_{b_h} \text{loss}_1^T & \cdots & \nabla_{b_h} \text{loss}_h^T \end{bmatrix}$$

4.0.2 Task 5 (5 points): Based on the previous implementation complete the following functions and train a 2 layers perceptron.

```
In [13]: def relu(x):
    """
    input:
        x (np.array)

    returns:
        x (np.array)
    """
    aux = x.reshape(-1)
    result = np.zeros(len(aux))
    for i in range(len(result)):
        result[i] = max(0, aux[i])
    return result.reshape(-1, 1)

def d_relu(x):
    """Computes the derivative of the relu
    input:
        x (np.array)

    returns:
        x (np.array)
    """
    aux = x.reshape(-1)
    result = np.zeros(len(aux))
```

```

for i in range(len(result)):
    if aux[i] > 0:
        result[i] = 1
    else:
        result[i] = 0
return result.reshape(-1,1)

def forward_two_layers(Wo, bo, Wh, bh, x):
    """Forward pass of a two layer perceptron with relu activation
    input:
        Wh (np.array): (INPUT_SHAPE, HIDDEN_SHAPE) The weight matrix of the
                        hidden layer
        Wo (np.array): (HIDDEN_SHAPE, N_CLASSES) The weight matrix of the
                        output layer
        bh (np.array): (HIDDEN_SHAPE, 1) The bias matrix of the hidden layer
        bo (np.array): (N_CLASSES, 1) The bias matrix of the output layer
        x (np.array): (INPUT_SHAPE, 1) The input of the perceptron

    returns:
        softmaxed (np.array): (N_CLASSES, 1) the output of the network after
                                final activation
        hidden (np.array): (HIDDEN_SHAPE, 1) the output of the hidden layer
                                after activation
        out (np.array): (N_CLASSES, 1) the output of the network before
                                final activation
    """

    hidden = relu(Wh.T @ x + bh) # h
    out = Wo.T @ hidden + bo     # y_hat
    softmaxed = softmax(out)     # Softmax(y_hat)

    return softmaxed, hidden, out

def compute_grads_two_layers(hidden, softmaxed, Wo, x, y):
    """Forward pass of a two layer perceptron with relu activation
    input:
        hidden (np.array): (HIDDEN_SHAPE, 1) the output of the hidden layer
                                after activation
        softmaxed (np.array): (N_CLASSES, 1) the output of the network after
                                final activation
        Wo (np.array): (HIDDEN_SHAPE, N_CLASSES) The weight matrix of the
                                output layer
        x (np.array): (INPUT_SHAPE, 1) The input of the perceptron
        y (np.array): (N_CLASSES, 1) Ground truth class

    returns:
        d_Wo (np.array): (HIDDEN_SHAPE, N_CLASSES) Gradient with respect
                                to the weight matrix of the output layer

```



```

        d_bo (np.array): (N_CLASSES, 1) Gradient with respect to the bias matrix
        of the output layer
        d_Wh (np.array): (INPUT_SHAPE, HIDDEN_SHAPE) Gradient with respect to the
        weight matrix of the hidden layer
        d_bh (np.array): (HIDDEN_SHAPE, 1) Gradient with respect to the bias
        matrix of the hidden layer
    """
    d_Wo, d_bo = compute_grads_one_layer(softmaxed, hidden, y)

    d_bh = np.zeros(hidden.shape)
    for i in range(d_bh.shape[0]):
        aux = d_relu(hidden)
        d_bh[i,0] = aux[i,0] * ( Wo[i,:] @ (softmaxed - y).reshape(-1) )

    d_Wh = np.outer(x, d_bh)

    return d_Wo, d_bo, d_Wh, d_bh

def train_two_layer(X_train, Y_train, X_test, Y_test, lr, n_hidden,
                    n_it=1000, test_freq=10, random_seed=42):

    INPUT_SHAPE = X_train.shape[1]
    N_CLASSES = Y_train.shape[1]

    # Initialise metrics lists
    loss = []
    acc = []
    test_acc = []
    approx_grads = []
    grads = []

    # Initialisation of the weights
    np.random.seed(random_seed)
    bh = rand.normal(size=(n_hidden, 1))
    Wh = rand.normal(size=(INPUT_SHAPE, n_hidden))
    bo = rand.normal(size=(N_CLASSES, 1))
    Wo = rand.normal(size=(n_hidden, N_CLASSES))

    # Shuffling data
    indexes = rand.randint(X_train.shape[0], size=n_it)
    # training loop
    for it, i in enumerate(indexes):
        x = X_train[i,:].reshape(-1,1)
        y = Y_train[i,:].reshape(-1,1)

        # Forward passs
        softmaxed, hidden, out = forward_two_layers(Wo, bo, Wh, bh, x)

```

```

# Back propagation
d_Wo, d_bo, d_Wh, d_bh = compute_grads_two_layers(hidden, softmaxed,
                                                    Wo, x, y)

bh -= lr * d_bh
Wh -= lr * d_Wh
bo -= lr * d_bo
Wo -= lr * d_Wo

# Metrics recording
loss.append(compute_loss(softmaxed, y))
acc.append(np.argmax(softmaxed) == np.argmax(y))

# Test loop
if it % test_freq == 0:
    acc_temp = []
    for i in range(X_test.shape[0]):
        x = X_train[i,:].reshape(-1,1)
        y = Y_train[i,:].reshape(-1,1)
        softmaxed, hidden, out = forward_two_layers(Wo, bo, Wh, bh, x)
        acc_temp.append(np.argmax(softmaxed) == np.argmax(y))

    test_acc.append(np.mean(acc_temp))

return Wo, bo, Wh, bh, loss, acc, test_acc

```

```

In [14]: Wo, bo, Wh, bh, loss, acc, test_acc = train_two_layer(X_train, Y_train,
                                                                X_test, Y_test, 0.1, 16)

```

```

iteration = [t+1 for t in range(1000)]

```

```

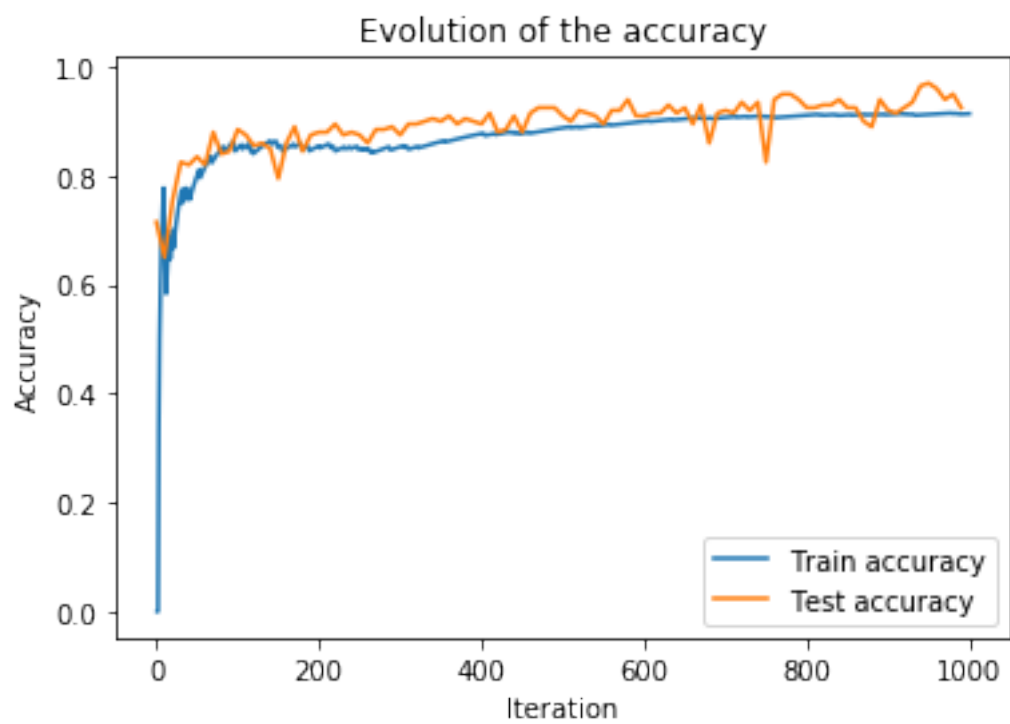

# evolution of loss function
plt.plot(iteration, loss)
plt.title("Evolution of the Loss function")
plt.xlabel("Iteration")
plt.ylabel("Loss")
plt.show()

```


```

# Training data accuracy
freq = 10
acc_true = [a/b for a,b in zip(np.array(acc).cumsum(), range(1,len(acc)+1))]
plt.plot(iteration, acc_true, label="Train accuracy")
plt.plot([t*freq for t in range(100)], test_acc, label="Test accuracy")
plt.xlabel("Iteration")
plt.ylabel("Accuracy")
plt.title("Evolution of the accuracy")
plt.legend()
plt.show()

```



```
In [15]: plot_decision(X_test, Y_test, lambda x: forward_two_layers(Wo, bo, Wh, bh, x)[0])
```




4.0.3 Task 7 (3 points): The non linearity of the decision function is conditioned by the number of units, visualize this effect. What can you comment on the smoothness of the boundary ? (Be careful, to visualize this you need to properly train the networks)

To better visualize this let's generate new data.

```
In [16]: X, Y = make_circles(n_samples=500, noise=.2, factor=0.4)
plt.figure(figsize=(10,10))
plt.scatter(X[:,0], X[:,1], c=Y)
plt.show()
Y = OneHotEncoder().fit_transform(Y.reshape(-1,1)).toarray()

X_train, X_test, Y_train, Y_test = train_test_split(X, Y, random_state=42,
                                                    test_size=0.4)
```



Answer:


We can see below the effect of the number of units of perceptrons in the hidden layer to the decision function.

According to the results below, we can clearly realize that the boundary of the decision function become more well defined and smoother with the increase of the units.


Using 2 units, the decision function looks linear. Using 5 units, the boundary is barely circular, but it is not well defined. With 10, 15, 20 and 50 units, the circular characteristic of the decision function is noticed, but they have different degrees of smoothness. And the more we use units of perceptrons in the hidden layer, the smoother the boundary becomes.

```
In [17]: values = [2, 5, 10, 15, 20, 50]
         for units_hidden in values:
             print("UNITS =", units_hidden)
             Wo, bo, Wh, bh, loss, acc, test_acc = train_two_layer(X_train, Y_train, X_test,
                                                                    Y_test, 0.1, units_hidden)
             plot_decision(X_test, Y_test, lambda x: forward_two_layers(Wo, bo, Wh, bh, x)[0])
```


UNITS = 2




UNITS = 5




UNITS = 10




UNITS = 15



UNITS = 20



UNITS = 50



4.0.4 Task 8 (2 points): To have a smooth boundary we need to change the activation. Change the activation from `relu` to `tanh`, and visualize the result.

If we use the activation to `tanh`, we have that:

$$\tanh(x) = \frac{e^{2x} - 1}{e^{2x} + 1}$$

$$\tanh'(x) = 1 - \tanh^2(x)$$

$$\begin{aligned} h_i &= \tanh\left(W_{h(i,\cdot)}^T x + b_{h(i)}\right) \\ &= \tanh\left(\sum_k w_{h(k,i)} x_k + b_{h(i)}\right) \end{aligned}$$

$$\begin{aligned} \frac{\partial h_i}{\partial W_{h(u,v)}} &= \begin{cases} 0 & , \quad i \neq v \\ x_u \tanh'\left(\sum_k w_{h(k,i)} x_k + b_{h(i)}\right) & , \quad i = v \end{cases} \\ &= \begin{cases} 0 & , \quad i \neq v \\ x_u(1 - h_i^2) & , \quad i = v \end{cases} \end{aligned}$$

$$\frac{\partial h_i}{\partial b_{h(v)}} = \begin{cases} 0 & , \quad i \neq v \\ 1 - h_i^2 & , \quad i = v \end{cases}$$

$$\frac{\partial \text{loss}}{\partial W_{h(u,v)}} = x_u(1 - h_v^2)W_{0(v,\cdot)} [\text{Softmax}(\hat{y}) - y]$$

$$\frac{\partial \text{loss}}{\partial b_{h(v)}} = (1 - h_v^2)W_{0(v,\cdot)} [\text{Softmax}(\hat{y}) - y]$$

```
In [18]: def tanh(x):
         """
         input:
             x (np.array)

         returns:
             x (np.array)
         """
         aux = x.reshape(-1)
         result = (np.exp(2*aux) - 1)/(np.exp(2*aux) + 1)
         return result.reshape(-1,1)

def d_tanh_hidden(hidden):
    """Computes the derivative of the tanh
    input:
        x (np.array)

    returns:
        x (np.array)
    """
    return 1 - hidden**2

def forward_two_layers(Wo, bo, Wh, bh, x):
    """Forward pass of a two layer perceptron with tanh activation
```

```

    input:
        Wh (np.array): (INPUT_SHAPE, HIDDEN_SHAPE) The weight matrix of the
            hidden layer
        Wo (np.array): (HIDDEN_SHAPE, N_CLASSES) The weight matrix of the
            output layer
        bh (np.array): (HIDDEN_SHAPE, 1) The bias matrix of the hidden layer
        bo (np.array): (N_CLASSES, 1) The bias matrix of the output layer
        x (np.array): (INPUT_SHAPE, 1) The input of the perceptron

    returns:
        softmaxed (np.array): (N_CLASSES, 1) the output of the network after
            final activation
        hidden (np.array): (HIDDEN_SHAPE, 1) the output of the hidden layer
            after activation
        out (np.array): (N_CLASSES, 1) the output of the network before
            final activation
    """

    hidden = tanh(Wh.T @ x + bh) # h
    out = Wo.T @ hidden + bo     # y_hat
    softmaxed = softmax(out)     # Softmax(y_hat)

    return softmaxed, hidden, out

def compute_grads_two_layers(hidden, softmaxed, Wo, x, y):
    """Forward pass of a two layer perceptron with tanh activation
    input:
        hidden (np.array): (HIDDEN_SHAPE, 1) the output of the hidden layer
            after activation
        softmaxed (np.array): (N_CLASSES, 1) the output of the network after
            final activation
        Wo (np.array): (HIDDEN_SHAPE, N_CLASSES) The weight matrix of the
            output layer
        x (np.array): (INPUT_SHAPE, 1) The input of the perceptron
        y (np.array): (N_CLASSES, 1) Ground truth class

    returns:
        d_Wo (np.array): (HIDDEN_SHAPE, N_CLASSES) Gradient with respect
            to the weight matrix of the output layer
        d_bo (np.array): (N_CLASSES, 1) Gradient with respect to the bias matrix
            of the output layer
        d_Wh (np.array): (INPUT_SHAPE, HIDDEN_SHAPE) Gradient with respect to the
            weight matrix of the hidden layer
        d_bh (np.array): (HIDDEN_SHAPE, 1) Gradient with respect to the bias
            matrix of the hidden layer
    """
    d_Wo, d_bo = compute_grads_one_layer(softmaxed, hidden, y)

```

```

    d_bh = np.zeros(hidden.shape)
    for i in range(d_bh.shape[0]):
        aux = d_tanh_hidden(hidden)
        d_bh[i,0] = aux[i,0] * ( Wo[i,:] @ (softmaxed - y).reshape(-1) )

    d_Wh = np.outer(x, d_bh)

    return d_Wo, d_bo, d_Wh, d_bh

In [19]: Wo, bo, Wh, bh, loss, acc, test_acc = train_two_layer(X_train, Y_train,
                                                                X_test, Y_test, 0.1, 16)

iteration = [t+1 for t in range(1000)]

# evolution of loss function
plt.plot(iteration, loss)
plt.title("Evolution of the Loss function")
plt.xlabel("Iteration")
plt.ylabel("Lrerror")
plt.show()

# Training data accuracy
freq = 10
acc_true = [a/b for a,b in zip(np.array(acc).cumsum(), range(1,len(acc)+1))]
plt.plot(iteration, acc_true, label="Train accuracy")
plt.plot([t*freq for t in range(100)], test_acc, label="Test accuracy")
plt.xlabel("Iteration")
plt.ylabel("Accuracy")
plt.title("Evolution of the accuracy")
plt.legend()
plt.show()

plot_decision(X_test, Y_test, lambda x: forward_two_layers(Wo, bo, Wh, bh, x)[0])

```

